

---

# COMP9318: Data Warehousing and Data Mining

— L2: Data Warehousing and OLAP —

---

- Why and What are Data Warehouses?

# Data Analysis Problems

---

- The same data found in many different systems
  - Example: customer data across different departments
  - The same concept is defined differently
- Heterogeneous sources
  - Relational DBMS, OnLine Transaction Processing (OLTP)
  - Unstructured data in files (e.g., MS Excel) and documents (e.g., MS Word)

# Data Analysis Problems (Cont'd)

---

- Data is suited for operational systems
  - Accounting, billing, etc.
  - Do not support analysis across business functions
- Data quality is bad
  - Missing data, imprecise data, different use of systems
- Data are “volatile”
  - Data deleted in operational systems (6months)
  - Data change over time – no historical information

# Solution: Data Warehouse

---

- Defined in many different ways, but not rigorously.
  - A decision support database that is maintained **separately** from the organization's operational database
  - Support **information processing** by providing a solid platform of consolidated, historical data for analysis.
- "A data warehouse is a subject-oriented, integrated, time-variant, and nonvolatile collection of data in support of management's decision-making process."—W. H. Inmon
- Data warehousing:
  - The process of constructing and using data warehouses

# Data Warehouse—Subject-Oriented

---

- Organized around major subjects, such as **customer, product, sales**.
- Focusing on the modeling and analysis of data for decision makers, not on daily operations or transaction processing.
- Provide **a simple and concise** view around particular subject issues by **excluding data that are not useful in the decision support process**.

# Data Warehouse—Integrated

---

- Constructed by integrating multiple, heterogeneous data sources
  - relational databases, flat files, on-line transaction records
- Data cleaning and data integration techniques are applied.
  - Ensure consistency in naming conventions, encoding structures, attribute measures, etc. among different data sources
    - E.g., Hotel price: currency, tax, breakfast covered, etc.
  - When data is moved to the warehouse, it is converted.

# Data Warehouse—Time Variant

---

- The time horizon for the data warehouse is significantly longer than that of operational systems.
  - Operational database: current value data.
  - Data warehouse data: provide information from a historical perspective (e.g., past 5-10 years)
- Every key structure in the data warehouse
  - Contains an element of time, explicitly or implicitly
  - But the key of operational data may or may not contain “time element”.



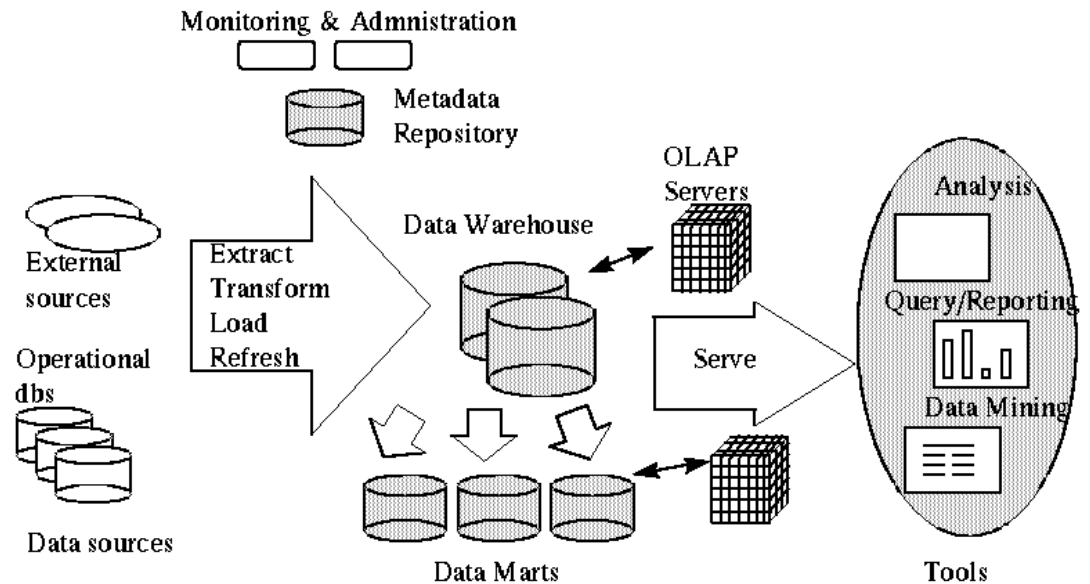
# Data Warehouse—Non-Volatile

---

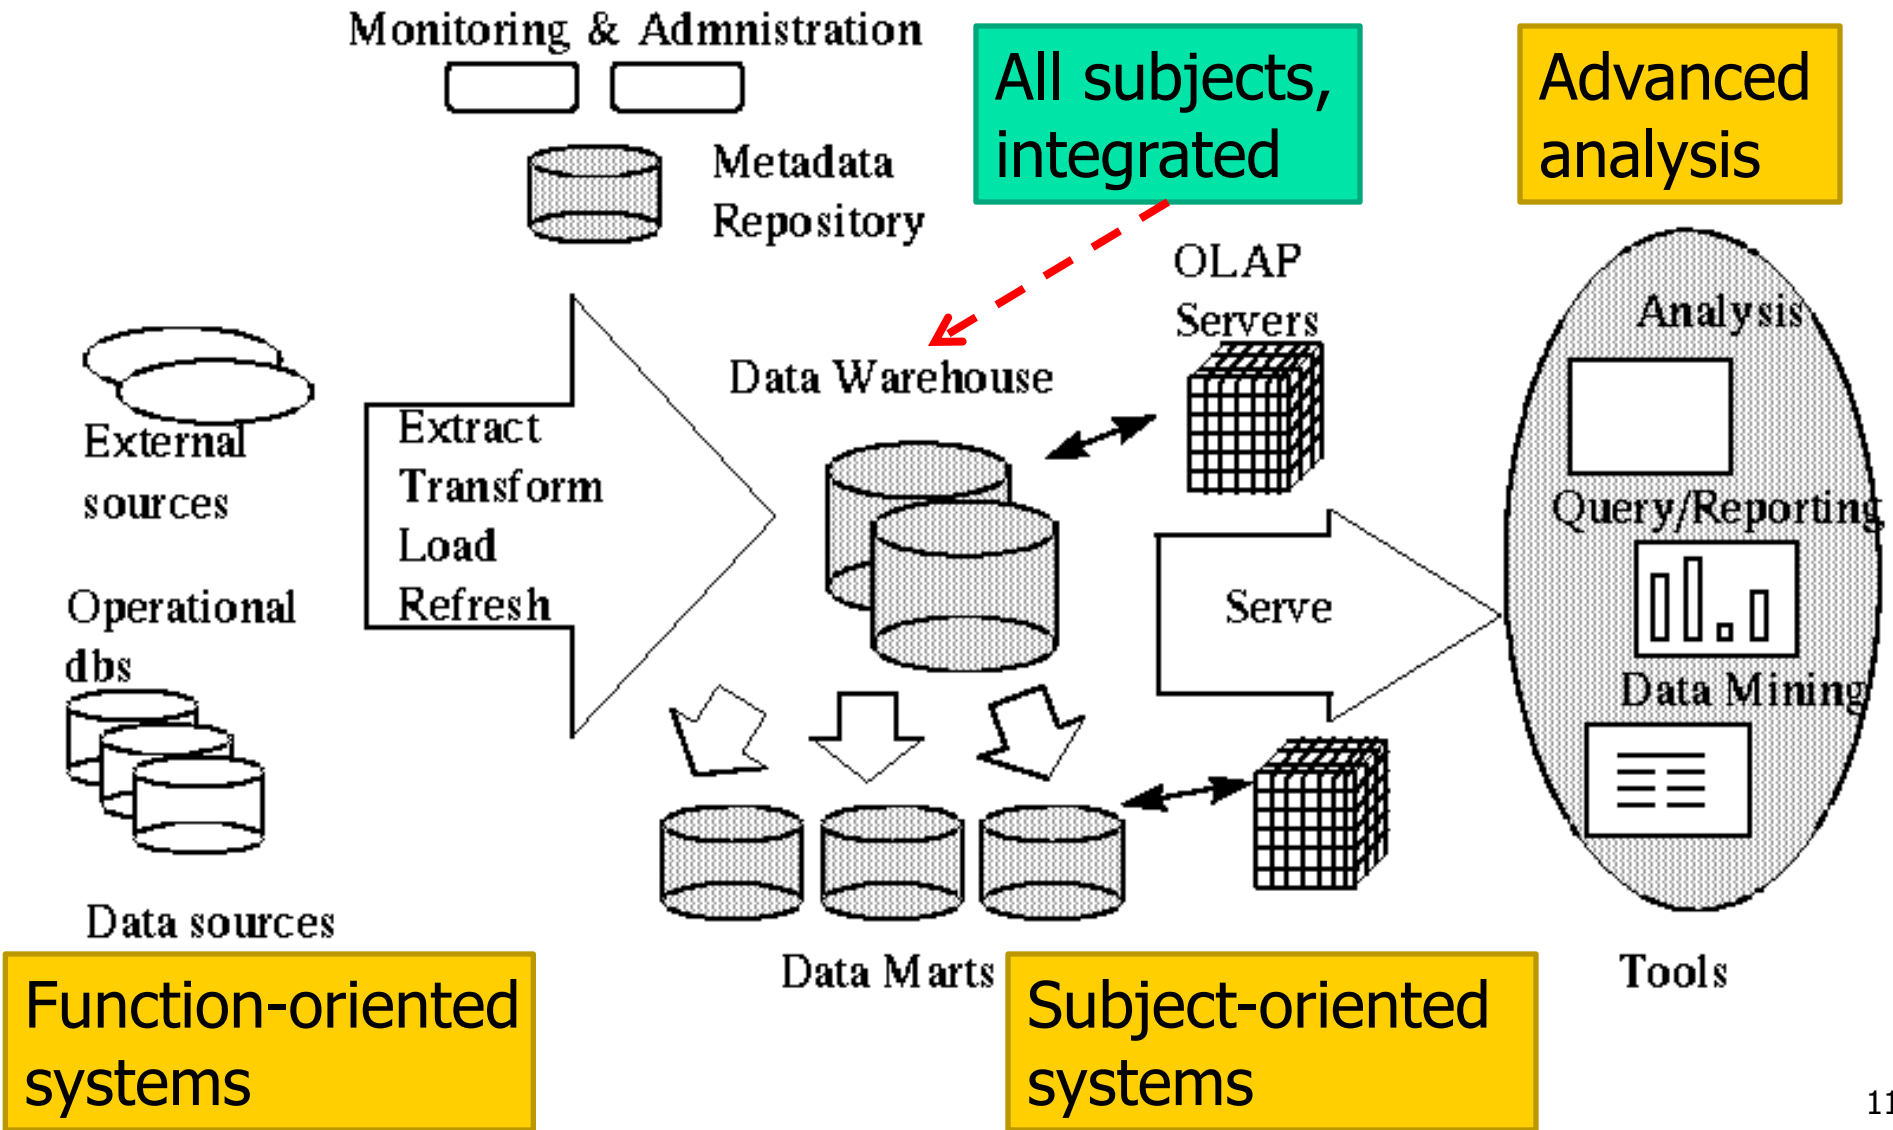
1. A **physically separate store** of data transformed from the operational environment.
2. Operational **update of data does not occur** in the data warehouse environment.
  - Does not require transaction processing, recovery, and concurrency control mechanisms
  - Requires only two operations in data accessing:
    - *initial loading of data* and *access of data*.

# Data Warehouse Architecture

- Extract data from operational data sources
  - clean, transform
- Bulk load/refresh
  - warehouse is offline
- OLAP-server provides multidimensional view
- **Multidimensional-olap**  
(Essbase, oracle express)
- **Relational-olap**  
(Redbrick, Informix, Sybase, SQL server)



# Data Warehouse Architecture



# Why Separate Data Warehouse?

---

- High performance for both systems
  - DBMS—tuned for OLTP: access methods, indexing, concurrency control, recovery
  - Warehouse—tuned for OLAP: complex OLAP queries, multidimensional view, consolidation.
- Different functions and different data:
  - missing data: Decision support requires historical data which operational DBs do not typically maintain
  - data consolidation: DS requires consolidation (aggregation, summarization) of data from heterogeneous sources
  - data quality: different sources typically use inconsistent data representations, codes and formats which have to be reconciled

# Why OLAP Servers?

---

- Different workload:
  - OLTP (on-line transaction processing)
    - Major task of traditional relational DBMS
    - Day-to-day operations: purchasing, inventory, banking, manufacturing, payroll, registration, accounting, etc.
  - OLAP (on-line analytical processing)
    - Major task of data warehouse system
    - Data analysis and decision making
- Queries hard/infeasible for OLTP, e.g.,
  - Which **week** we have the largest sales?
  - Does the sales of **dairy products** increase over time?
  - Generate a **spread sheet** of total sales by state and by year.
- Difficult to represent these queries by using SQL ← Why?

# OLTP vs. OLAP

---

	<b>OLTP</b>	<b>OLAP</b>
<b>users</b>	clerk, IT professional	knowledge worker
<b>function</b>	day to day operations	decision support
<b>DB design</b>	application-oriented	subject-oriented
<b>data</b>	current, up-to-date detailed, flat relational isolated	historical, summarized, multidimensional integrated, consolidated
<b>usage</b>	repetitive	ad-hoc
<b>access</b>	read/write index/hash on prim. key	lots of scans
<b>unit of work</b>	short, simple transaction	complex query
<b># records accessed</b>	tens	millions
<b>#users</b>	thousands	hundreds
<b>DB size</b>	100MB-GB	100GB-TB
<b>metric</b>	transaction throughput	query throughput, response

# Comparisons

	Databases	Data Warehouses
Purpose	Many purposes; Flexible and general	One purpose: Data analysis
Conceptual Model	ER	Multidimensional
Logical Model	(Normalized) Relational Model	(Denormalized) Star schema / Data cube/cuboids
Physical Model	Relational Tables	ROLAP: Relational tables MOLAP: Multidimensional arrays
Query Language	SQL (hard for analytical queries)	MDX (easier for analytical queries)
Query Processing	B+-tree/hash indexes, Multiple join optimization, Materialized views	Bitmap/Join indexes, Star join, Materialized data cube

---

- The Multidimensional Model



# The Multidimensional Model

---

- A data warehouse is based on a multidimensional data model which views data in the form of a **data cube**, which is a multidimensional generalization of 2D spread sheet.
- Key concepts:
  - **Facts**: the subject it models
    - Typically transactions in this course; other types includes snapshots, etc.
    - Measures: numbers that can be aggregated
    - Dimensions: context of the measure
  - Hierarchies:
    - Provide contexts of different granularities (aka. grains)
- Goals for dimensional modeling:
  - Surround facts with as much relevant context (dimensions) as possible ← **Why?**

# Supermarket Example

---

- Subject: analyze total sales and profits
- Fact: Each Sales **Transaction**
  - Measure: Dollars\_Sold, Amount\_Sold, Cost
  - Calculated Measure: Profit
- Dimensions:
  - Store
  - Product
  - Time

# Visualizing the Cubes

- A valid **instance** of the model is a data cube

The diagram illustrates the transformation of a data cube from a fact table to a sparse matrix representation. A yellow arrow points from the fact table on the left to the sparse matrix on the right.

**Fact Table (Left):**

total Sales		product			
		p1	p2	p3	p4
city	NY	\$454	-	-	\$925
	LA	\$468	\$800	-	-
	SD	\$296	-	\$240	-
	SF	\$652	-	\$540	\$745

**Sparse Matrix (Right):**

city	454			925
	468	800		
	296		240	
	652		540	745
	product			


**Concepts:** cell, fact (=non-empty cell), measure, dimensions

Q: How to generalize it to 3D?

# 3D Cube and Hierarchies

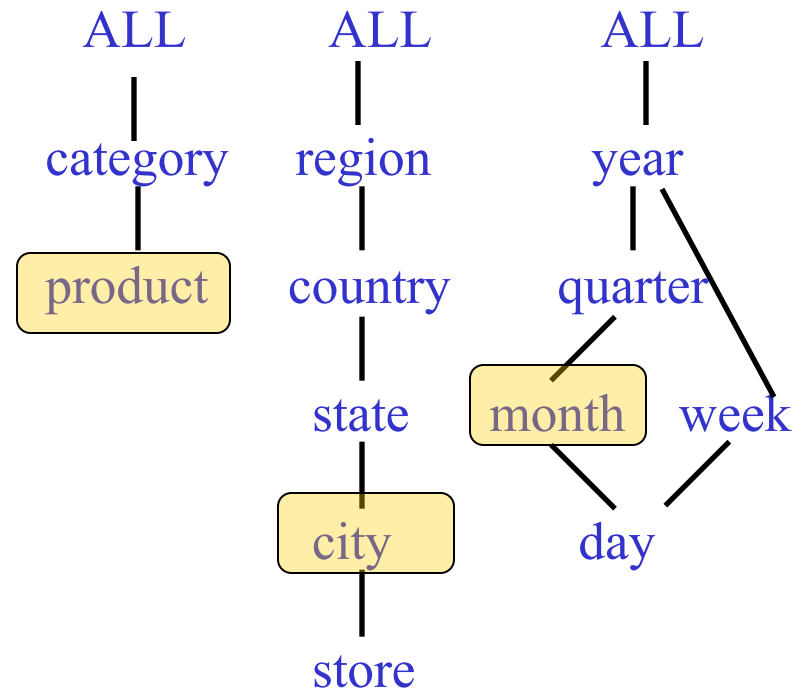
**Concepts:** hierarchy (a tree of dimension values), level

Sales of **book176** in **NY** in **Jan** can be found in this cell



## DIMENSIONS

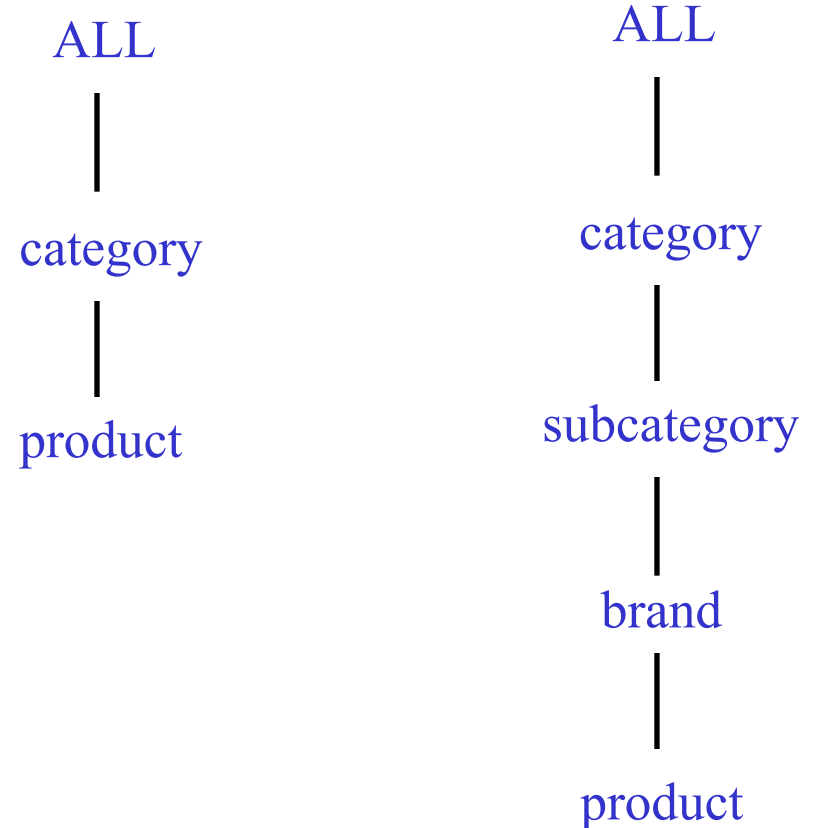

PRODUCT LOCATION TIME



# Hierarchies


**Concepts:** hierarchy (a tree of dimension values), level

## Which design is better? Why?



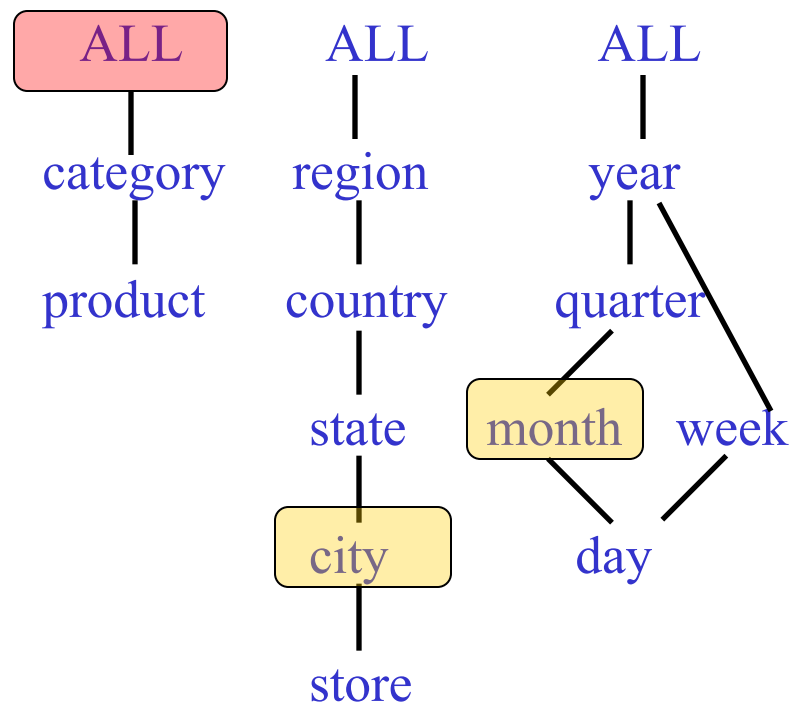
# The (city, moth) Cuboid

Sales of **ALL\_PROD** in **NY** in **Jan**




## DIMENSIONS

PRODUCT LOCATION TIME




Assume: no other non-ALL levels on all dimensions.

## All the Cuboids



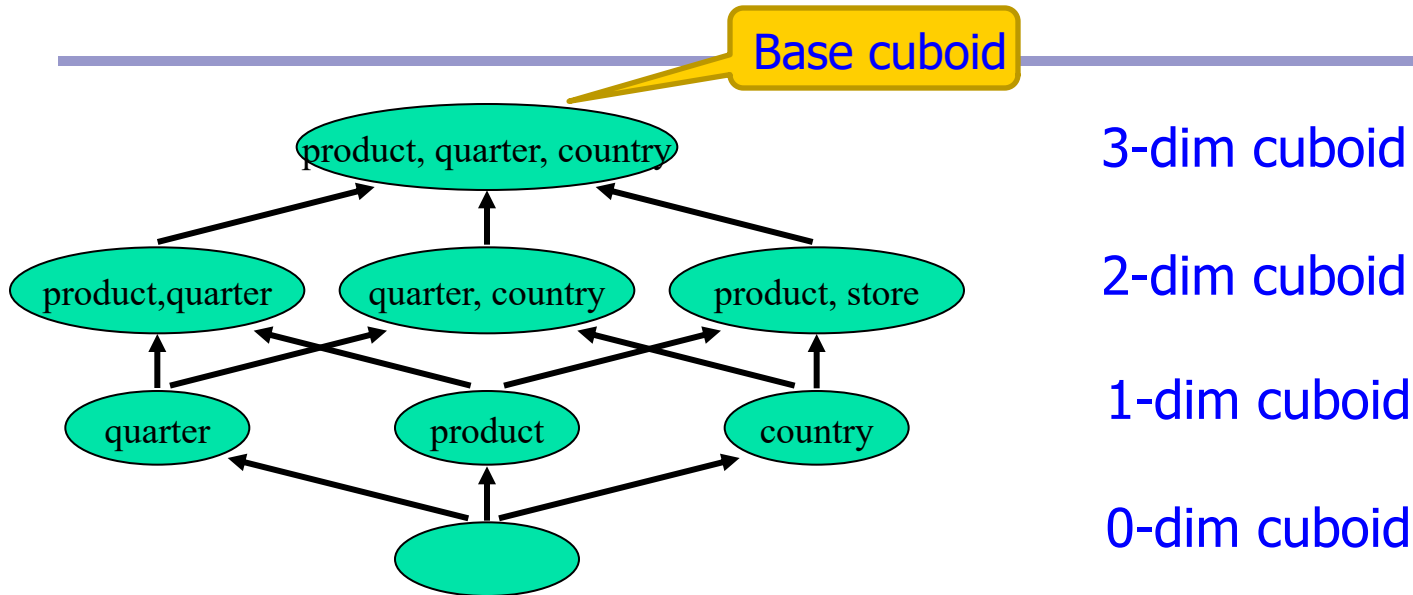
## All the Cuboids /2

Assume: no other non-ALL levels on all dimensions.





# Lattice of the cuboids



- n-dim cube can be represented as  $(D_1, D_2, \dots, D_d)$ , where  $D_i$  is the set of allowed values on the i-th dimension.
  - if  $D_i = L_i$  (a particular level), then  $D_i =$  all descendant dimension values of  $L_i$ .
  - ALL can be omitted and hence reduces the effective dimensionality
- A complete cube of d-dimensions consists of  $\prod_{i=1}^d (n_i + 1)$  cuboids, where  $n_i$  is the number of levels (excluding ALL) on i-th dimension.
  - They collectively form a lattice.

# Properties of Operations

---

- All operations are closed under the multidimensional model
  - i.e., both input and output of an operation is a cube
- So that they can be composed


Q: What's the analogy in the Relational Model?

# Common OLAP Operations

- **Roll-up**: move up the hierarchy

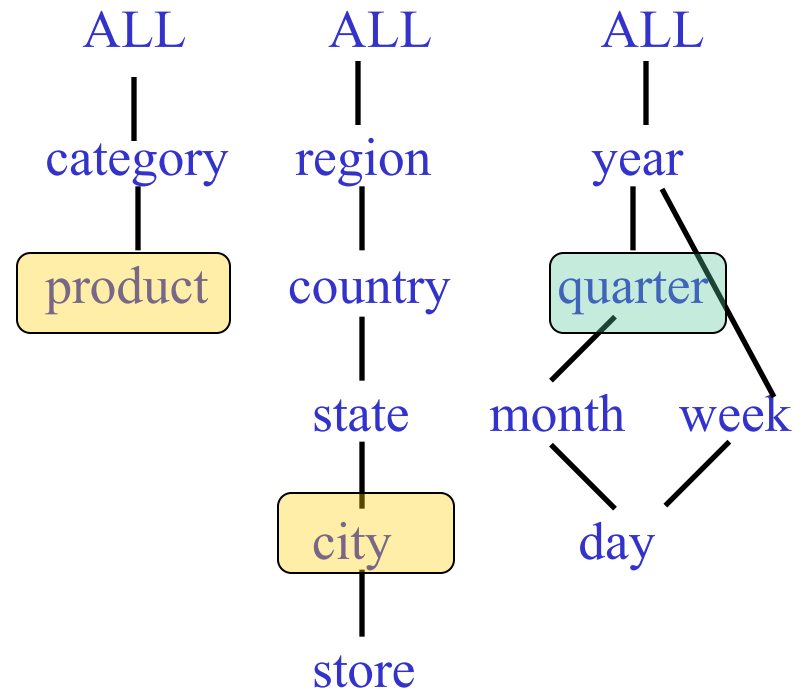
Q: what should be its value?



Sales of **book176** in **NY** in **Q1** here



## DIMENSIONS

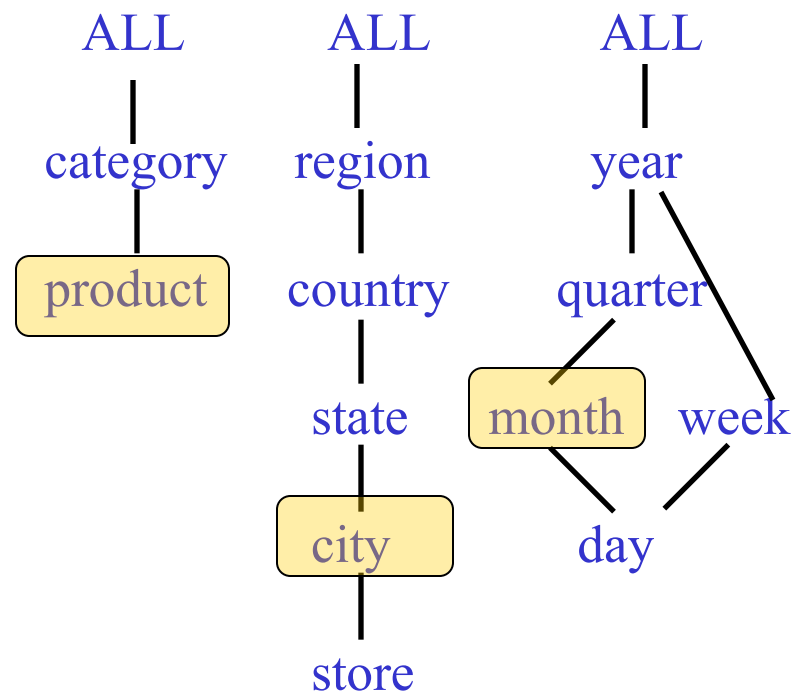
PRODUCT LOCATION TIME





# DIMENSIONS

## PRODUCT LOCATION TIME





# Common OLAP Operations

---

- **Drill-down**: move down the hierarchy
  - more fine-grained aggregation

# Slice and Dice Queries

- Slice and Dice: select and project on one or more dimension values





The output cube has smaller dimensionality than the input cube

# Pivoting

- Pivoting: aggregate on selected dimensions
  - usually 2 dims (cross-tabulation)

Sales (of **all products**) in  
**NY** in **Q1**  
=sum( **???** )



pivot on (city, month)

A 2D grid representing a pivot table. The vertical axis is labeled 'City' in blue, with 'NY' in red to its left. The horizontal axis is labeled 'month' in blue, with 'Q1', 'Q2', and 'Q3' in red below it. The grid contains several cells, some with ellipses (...) indicating data. A dashed blue line with an arrow points from the 'NY' label to the first row of the grid.

...		...
...	...	
...		...
...		...

# A Reflective Pause

---

- Let's review the definition of data cubes again.
- **Key message:**
  - Disentangle the “object” from its “representation” or “implementation”



# Modeling Exercise 1: Monthly Phone Service Billing

The screenshot displays the Telstra website interface. At the top, the Telstra logo is on the left, and a search bar is on the right. Below the logo, a navigation bar contains links: Telstra Home, Shop Online, My Account, Products & Services, Help Centre, and About Telstra. The main content area is divided into several sections. On the left, there's a 'My Account' section with a 'Login to My Account' button, and a 'Quick Links' section with a list of links: Check my WebMail, Pay my bill, Find a mobile, Move home, Activate Pre-Paid, Search FAQs, Switch to Telstra, Find a Telstra or T[life] shop, Discover Next G™, and Contact Us. Below this is a 'Latest News' section. The central part of the page features two large promotional banners. The left banner is for 'Whereis® Mobile' on the 'Next G™' network, showing a mobile phone and a map. The right banner is for 'Next G™ handset coverage questions', featuring a 'RECOMMENDED' seal and the 1800 888 888 hotline. Below these banners, the page is split into 'Personal' and 'Business' sections. The 'Personal' section includes links for 'Mobile', 'BigPond® Internet', 'Home Phone', and 'FOXTEL AUSTAR', along with a 'Bundle & Save' promotion. The 'Business' section includes links for 'Small & Medium Business' and 'Enterprise & Government', with a 'Login to Online Services' button. The bottom of the page features a green banner with the text 'Theme: analyze the income/revenue of Telstra'.

**Telstra**

Search

Telstra

Telstra Home | Shop Online | My Account | Products & Services | Help Centre | About Telstra

**My Account**

Login to My Account

**Quick Links**

- » Check my WebMail
- » Pay my bill
- » Find a mobile
- » Move home
- » Activate Pre-Paid
- » Search FAQs
- » Switch to Telstra
- » Find a Telstra or T[life] shop
- » Discover Next G™
- » Contact Us

**Latest News**

**Whereis® Mobile** - FREE to browse^ on your Next G™ mobile

For Next G™ handset coverage questions call the **1800 888 888** hotline

Find out more

**Personal**

Mobile | BigPond® Internet

Bundle & Save | Home Phone | FOXTEL AUSTAR

from Telstra

**Business**

Small & Medium Business | Enterprise & Government

Login to Online Services | Visit us today >

Theme: analyze the income/revenue of Telstra

# Solution

---

- FACT
- MEASURE
- DIMENSIONS



---

- The Logical Model

# Logical Models

---

- Two main approaches:
  - Using relational DB technology:
    - Star schema, Snowflake schema, Fact constellation
  - Using multidimensional technology:
    - Just as multidimensional data cube


# Universal Schema → Star Schema

- Many data warehouses adopt a star schema to represent the multidimensional model
- Each dimension is represented by a dimension-table
  - LOCATION (location\_key, store, street\_address, city, state, country, region)
  - dimension tables are not normalized
- Transactions are described through a fact-table
  - each tuple consists of a logical pointer to each of the dimension-tables (foreign-key) and a list of measures (e.g. sales \$\$\$)

## The universal schema for supermarket

Store	City	State	Prod	Brand	Category	\$Sold	#Sold	Cost
S136	Syd	NSW	76Ha	Nestle	Biscuit	40	10	18
S173	Melb	Vic	76Ha	Nestle	Biscuit	20	5	11

# The Star Schema



Think why:

- (1) Denormalized **once** from the universal schema
- (2) Controlled **redundancy**

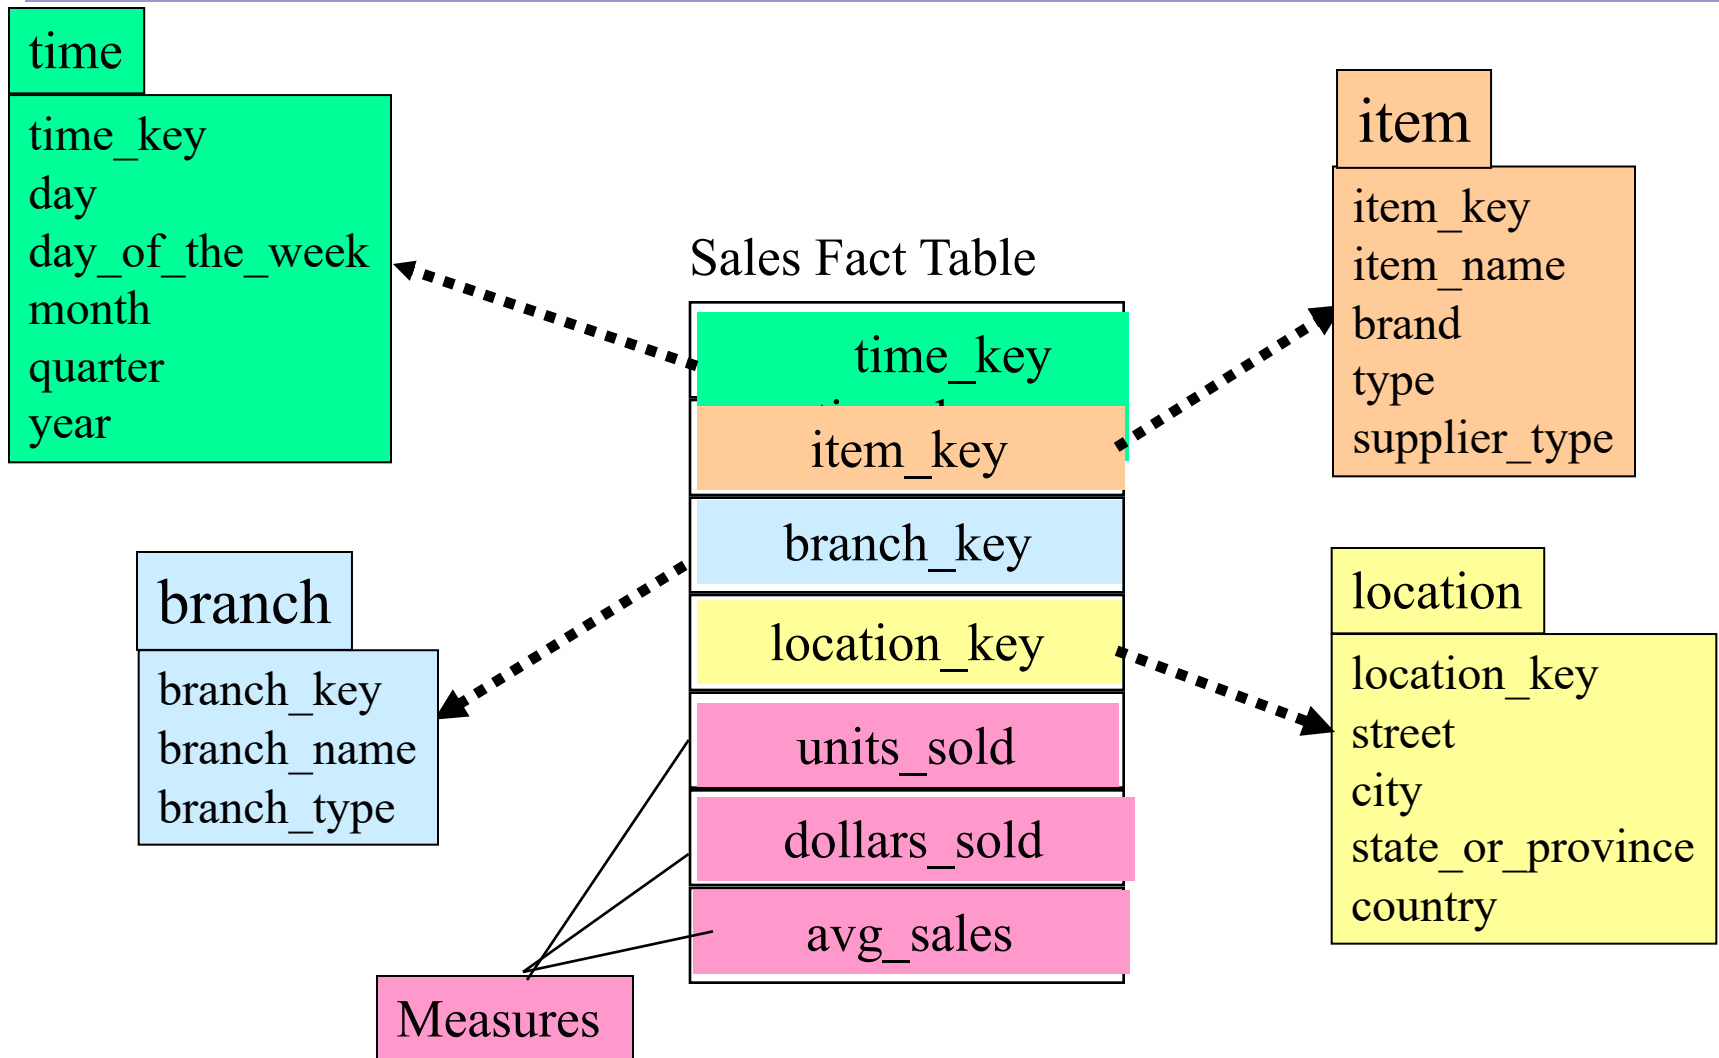
# Typical Models for Data Warehouses

---

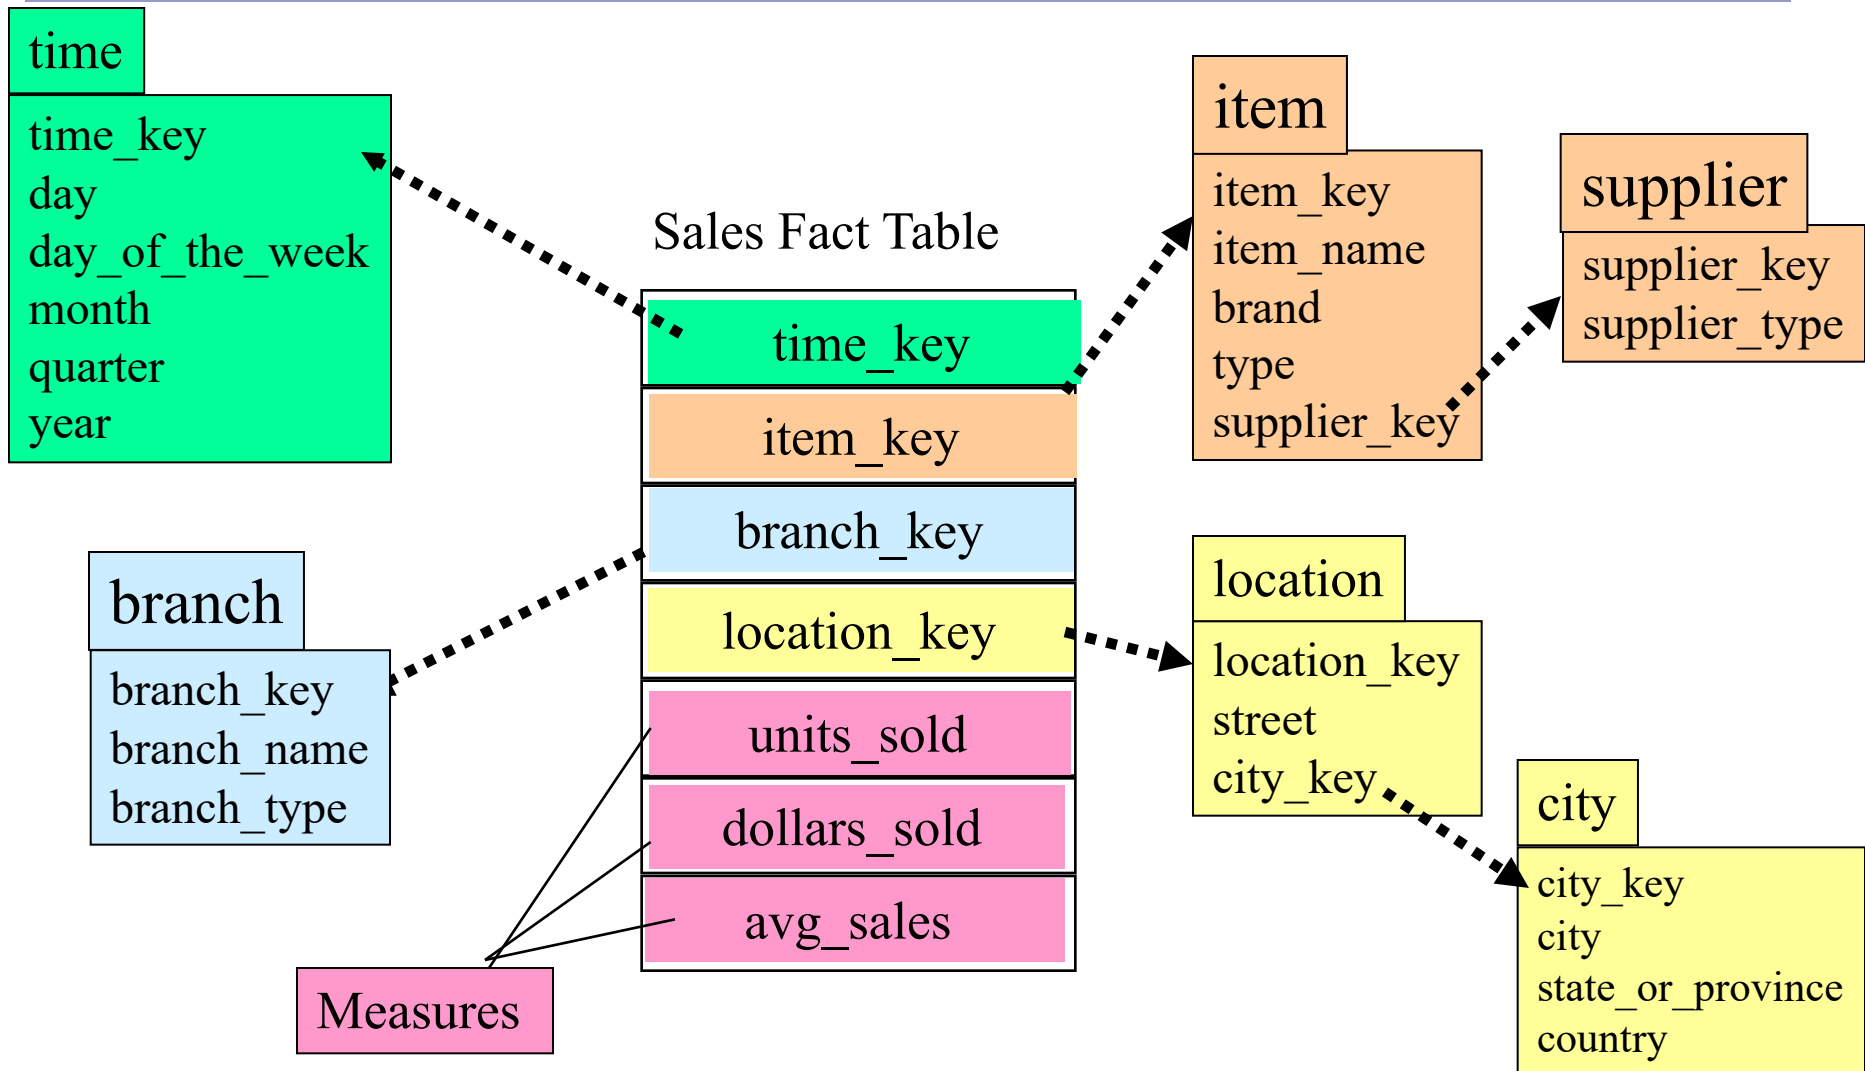
- Modeling data warehouses: dimensions & measures
  - Star schema: A fact table in the middle connected to a set of dimension tables
  - Snowflake schema: A refinement of star schema where some dimensional hierarchy is **normalized** into a set of smaller dimension tables, forming a shape similar to snowflake
  - Fact constellations: Multiple fact tables share dimension tables, viewed as a collection of stars, therefore called **galaxy schema** or fact constellation



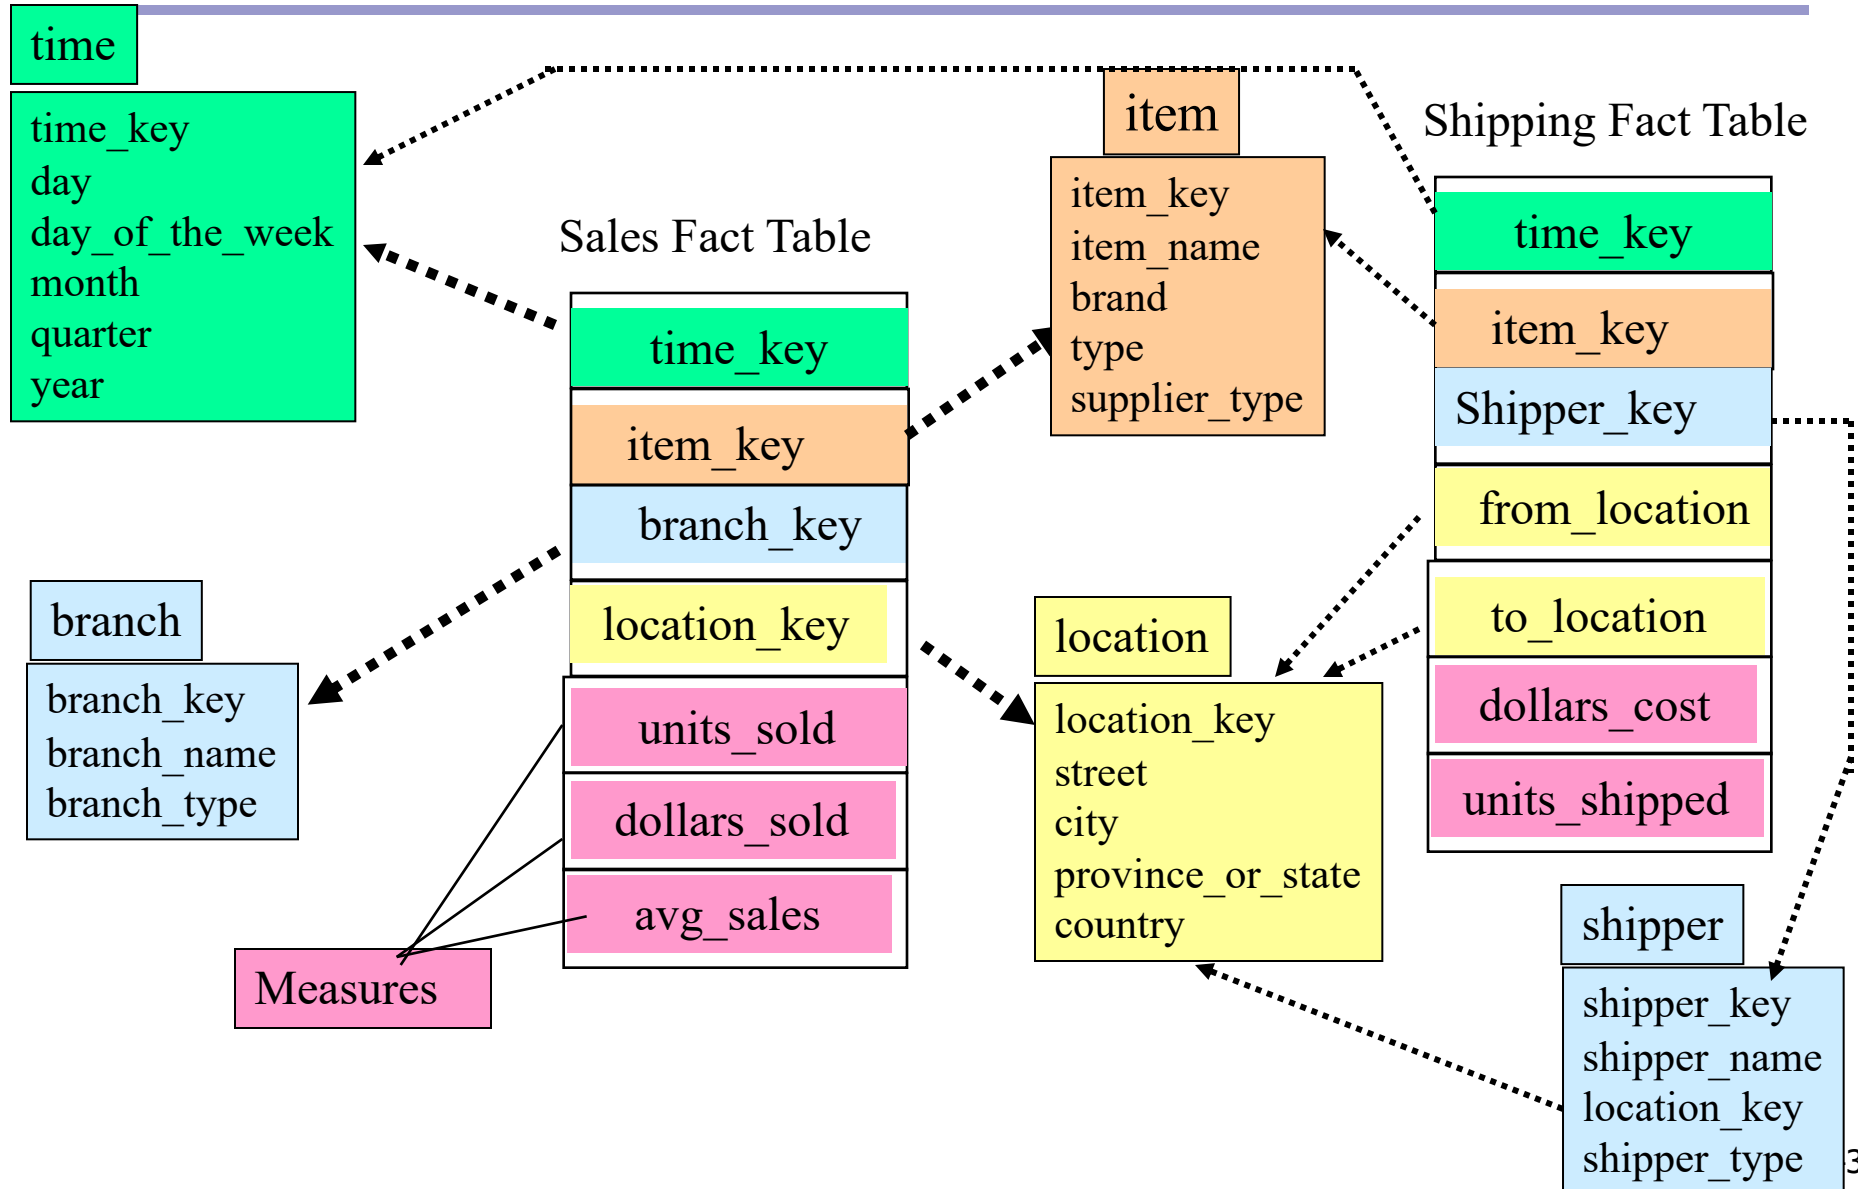
# Example of Star Schema



# Example of Snowflake Schema



# Example of Fact Constellation



# Advantages of Star Schema

---

- Facts and dimensions are clearly depicted
  - dimension tables are relatively static, data is loaded (append mostly) into fact table(s)
  - easy to comprehend (and write queries)

*“Find total sales per product-category in our stores in Europe”*

```
SELECT PRODUCT.category, SUM(SALES.amount)
FROM    SALES, PRODUCT, LOCATION
WHERE    SALES.product_key = PRODUCT.product_key
AND      SALES.location_key = LOCATION.location_key
AND      LOCATION.region="Europe"
GROUP BY PRODUCT.category
```

**Operations:** Slice (Loc.Region.Europe) + Pivot (Prod.category)

---

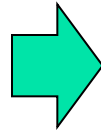
- Query Language

# Query Language

---

- Two approaches:
  - Using relational DB technology: SQL (with extensions such as CUBE/PIVOT/UNPIVOT)
  - Using multidimensional technology: MDX

```
SELECT PRODUCT.category,  
SUM(SALES.amount)  
FROM SALES, PRODUCT, LOCATION  
WHERE SALES.product_key =  
PRODUCT.product_key  
AND SALES.location_key =  
LOCATION.location_key  
AND LOCATION.region="Europe"  
GROUP BY PRODUCT.category
```



```
SELECT  
  {[PRODUCT].[category]} on ROWS,  
  {[MEASURES].[amount]} on COLUMNS  
FROM [SALES]  
WHERE ([LOCATION].[region].[Europe])
```

Operations: Slice (Loc.Region.Europe) + Pivot (Prod.category, Measures.amnt)

---

- Physical Model + Query Processing Techniques

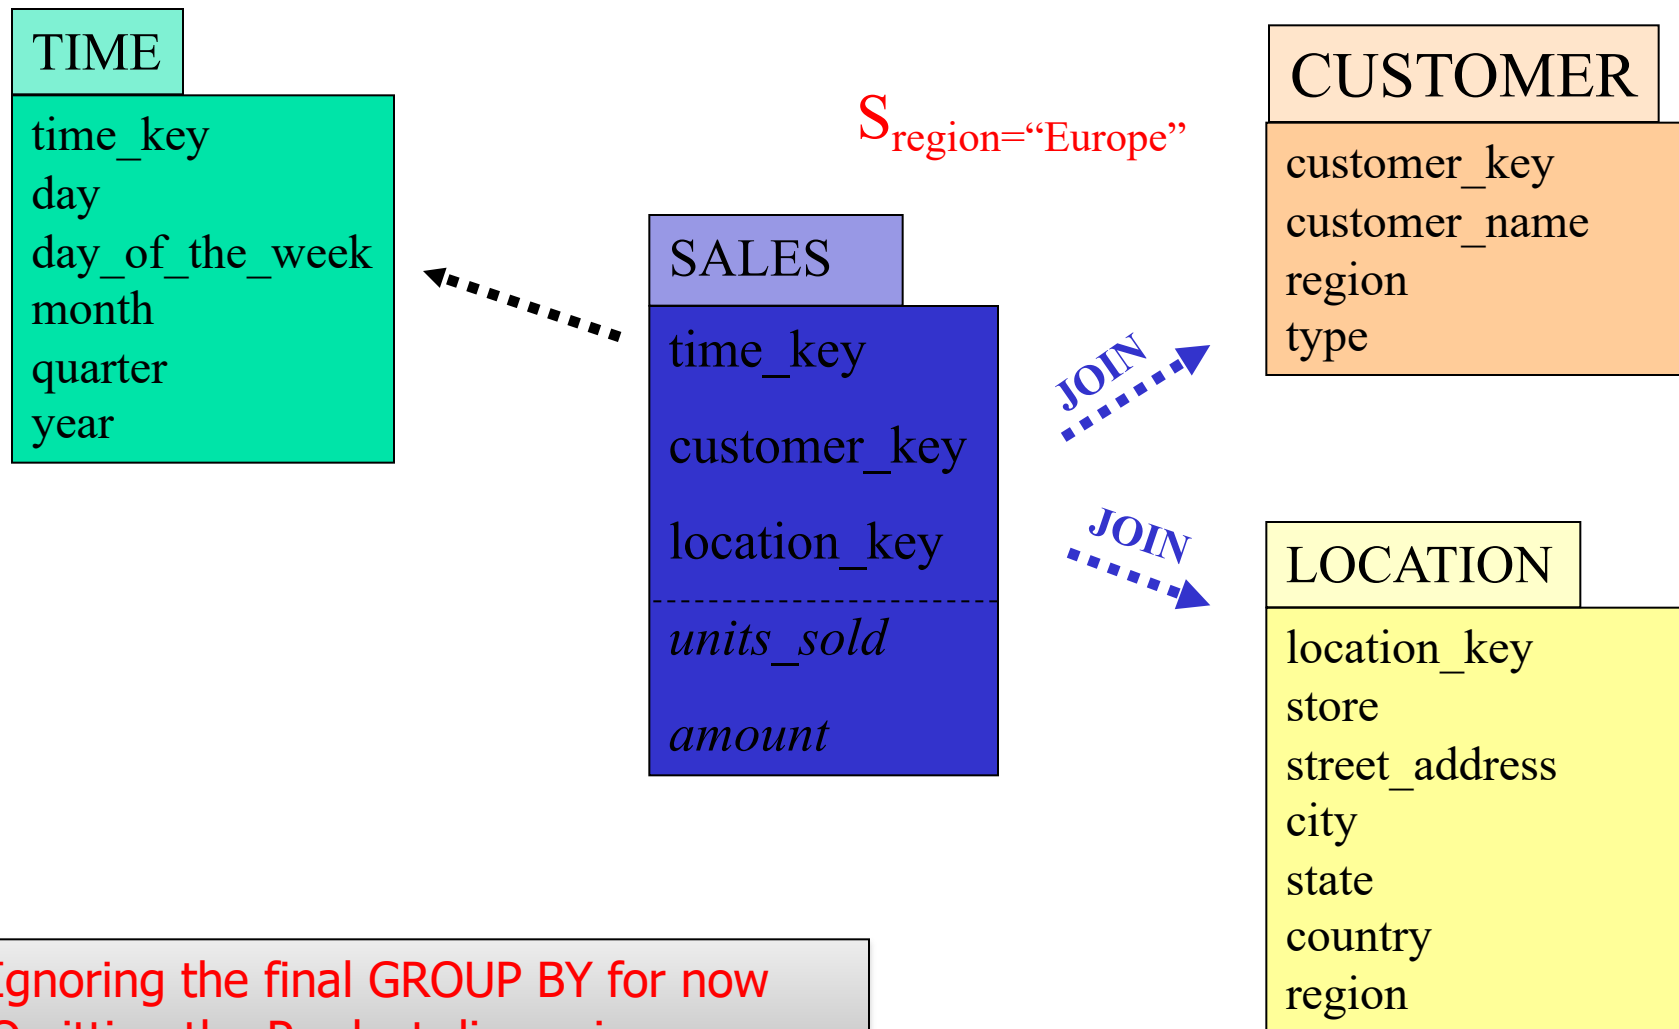
# Physical Model + Query Processing Techniques

---

- Two main approaches:
  - Using relational DB technology: ROLAP
  - Using multidimensional technology: MOLAP
- Hybrid: HOLAP
  - Base cuboid: ROLAP
  - Other cuboids: MOLAP



# Q1: Selection on low-cardinality attributes



- Ignoring the final GROUP BY for now
- Omitting the Product dimension

# Indexing OLAP Data: Bitmap Index

## (1) BI on dimension tables

- Index on an attribute (column) with low distinct values
- Each distinct values,  $v$ , is associated with a  $n$ -bit vector ( $n = \#rows$ )
  - The  $i$ -th bit is set if the  $i$ -th row of the table has the value  $v$  for the indexed column
- Multiple BIs can be efficiently combined to enable optimized scan of the table

### Custom

Cust	Region	Type
C1	Asia	Retail
C2	Europe	Dealer
C3	Asia	Dealer
C4	America	Retail
C5	Europe	Dealer

### BI on Customer.Region

$v$	bitmap
Asia	1 0 1 0 0
Europe	0 1 0 0 1
America	0 0 0 1 0

# Indexing OLAP Data: Bitmap Index /2

---

- (1) Bitmap join index (BI on Fact Table Joined with Dimension tables)
  - Conceptually, perform a join, map each dimension value to the bitmap of corresponding fact table rows.

-- ORACLE SYNTAX --

```
CREATE BITMAP INDEX sales_cust_region_bjix
ON   sales(customer.cust_region)
FROM sales, customer
WHERE sales.cust_id = customers.cust_id;
```

# Indexing OLAP Data: Bitmap Index /3

## Sales

time	customer	loc	Sale
101	C1	100	1
173	C1	200	2
208	C2	100	3
863	C3	200	5
991	C1	100	8
1001	C2	200	13
1966	C4	100	21
2017	C5	200	34

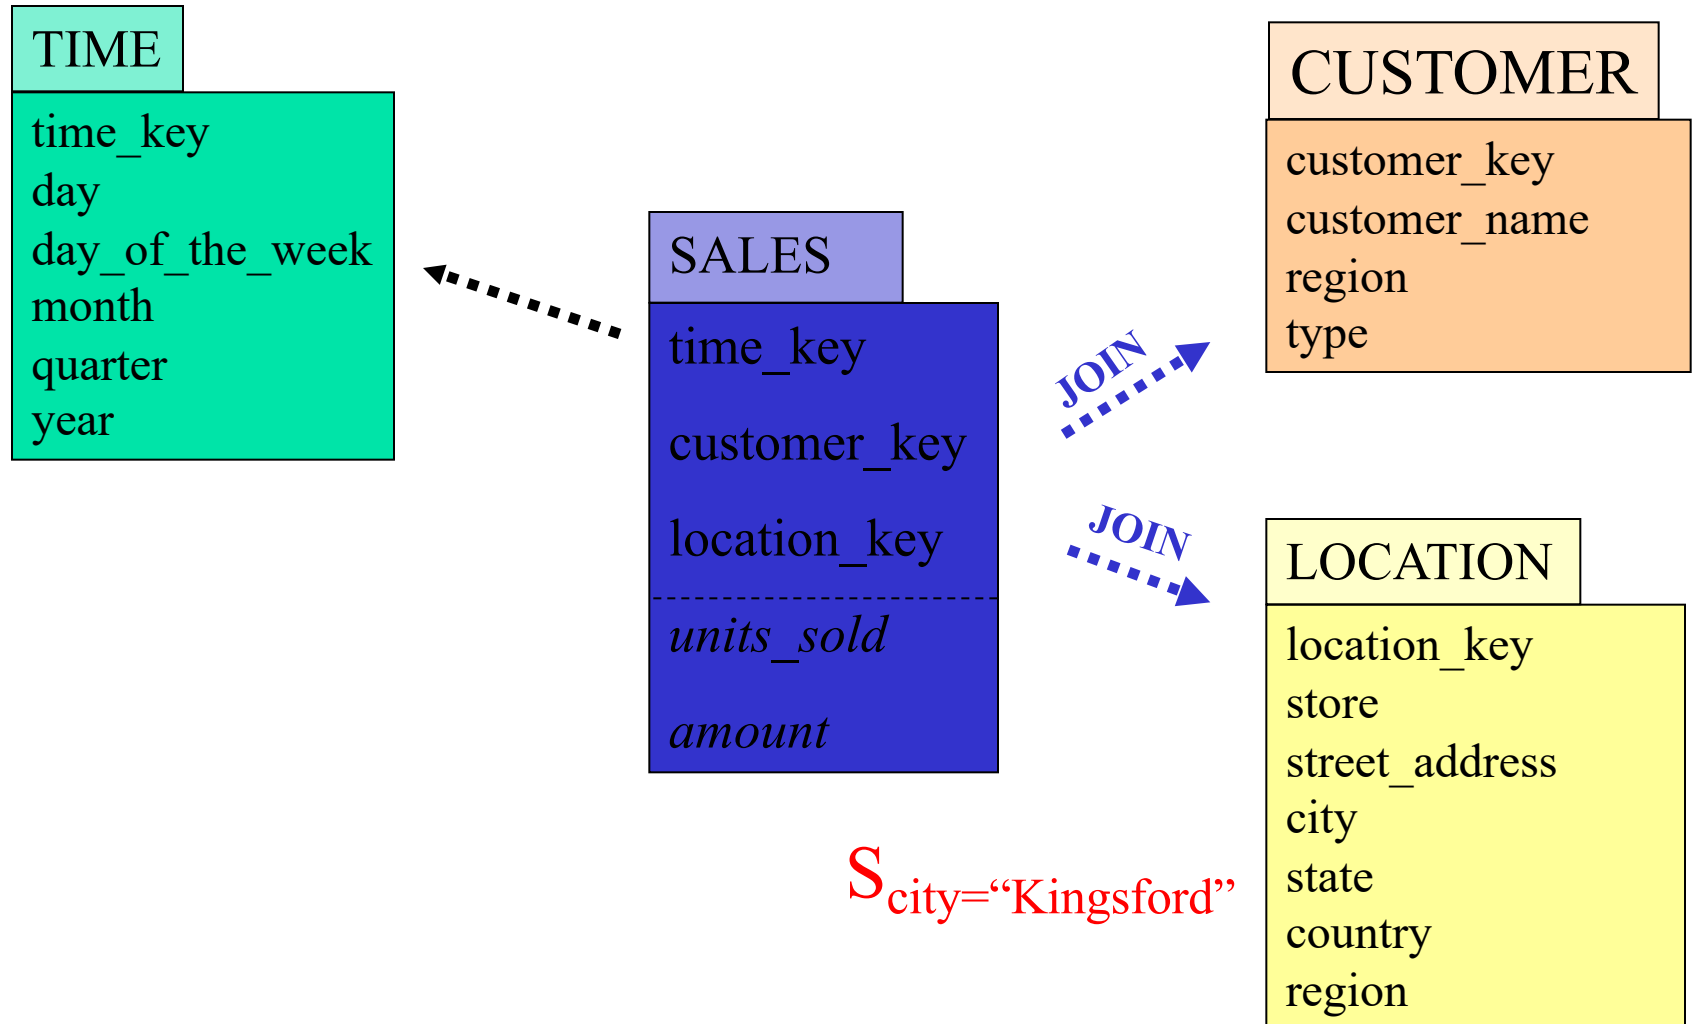
## Customer

Cust	Region	Type
C1	Asia	Retail
C2	Europe	Dealer
C3	Asia	Dealer
C4	America	Retail
C5	Europe	Dealer

## BI on Sales(Customer.Region)

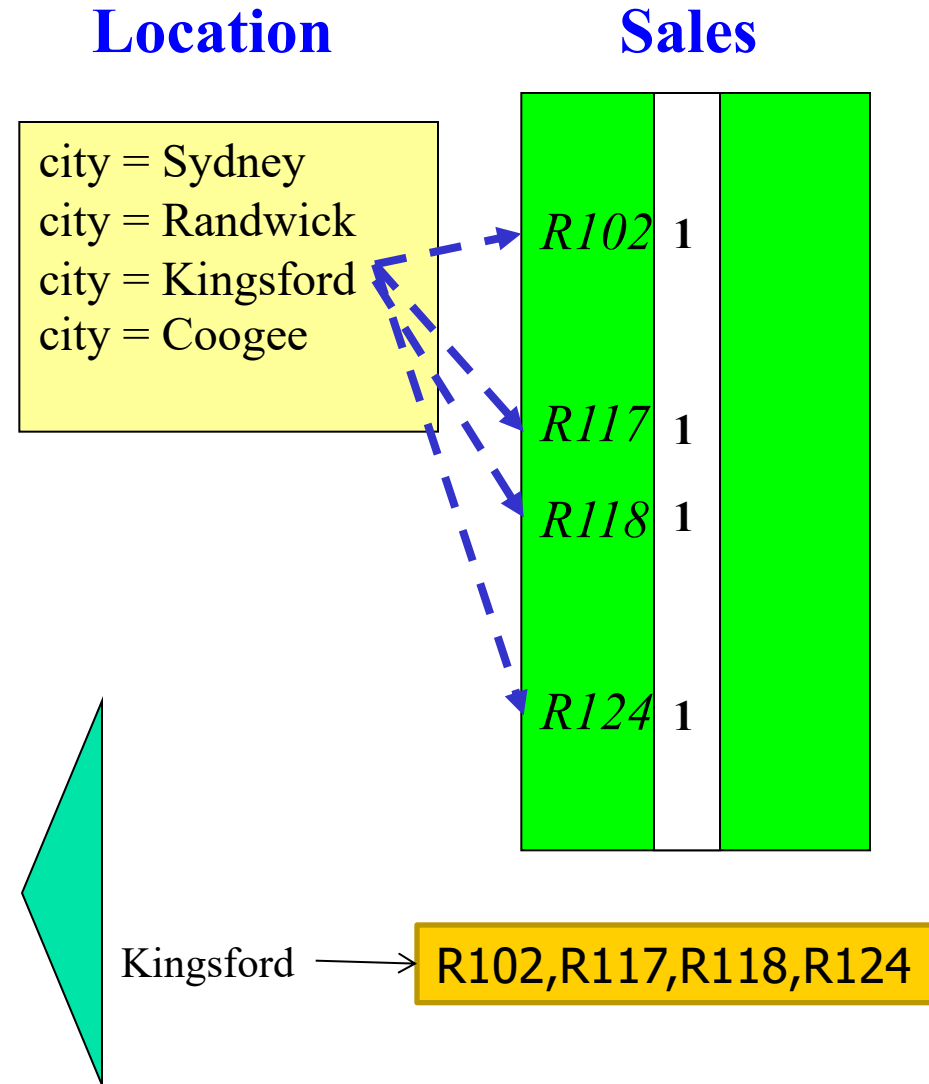
v	bitmap
Asia	11011000
Europe	00100101
America	00000010

# Q2: Selection on high-cardinality attributes

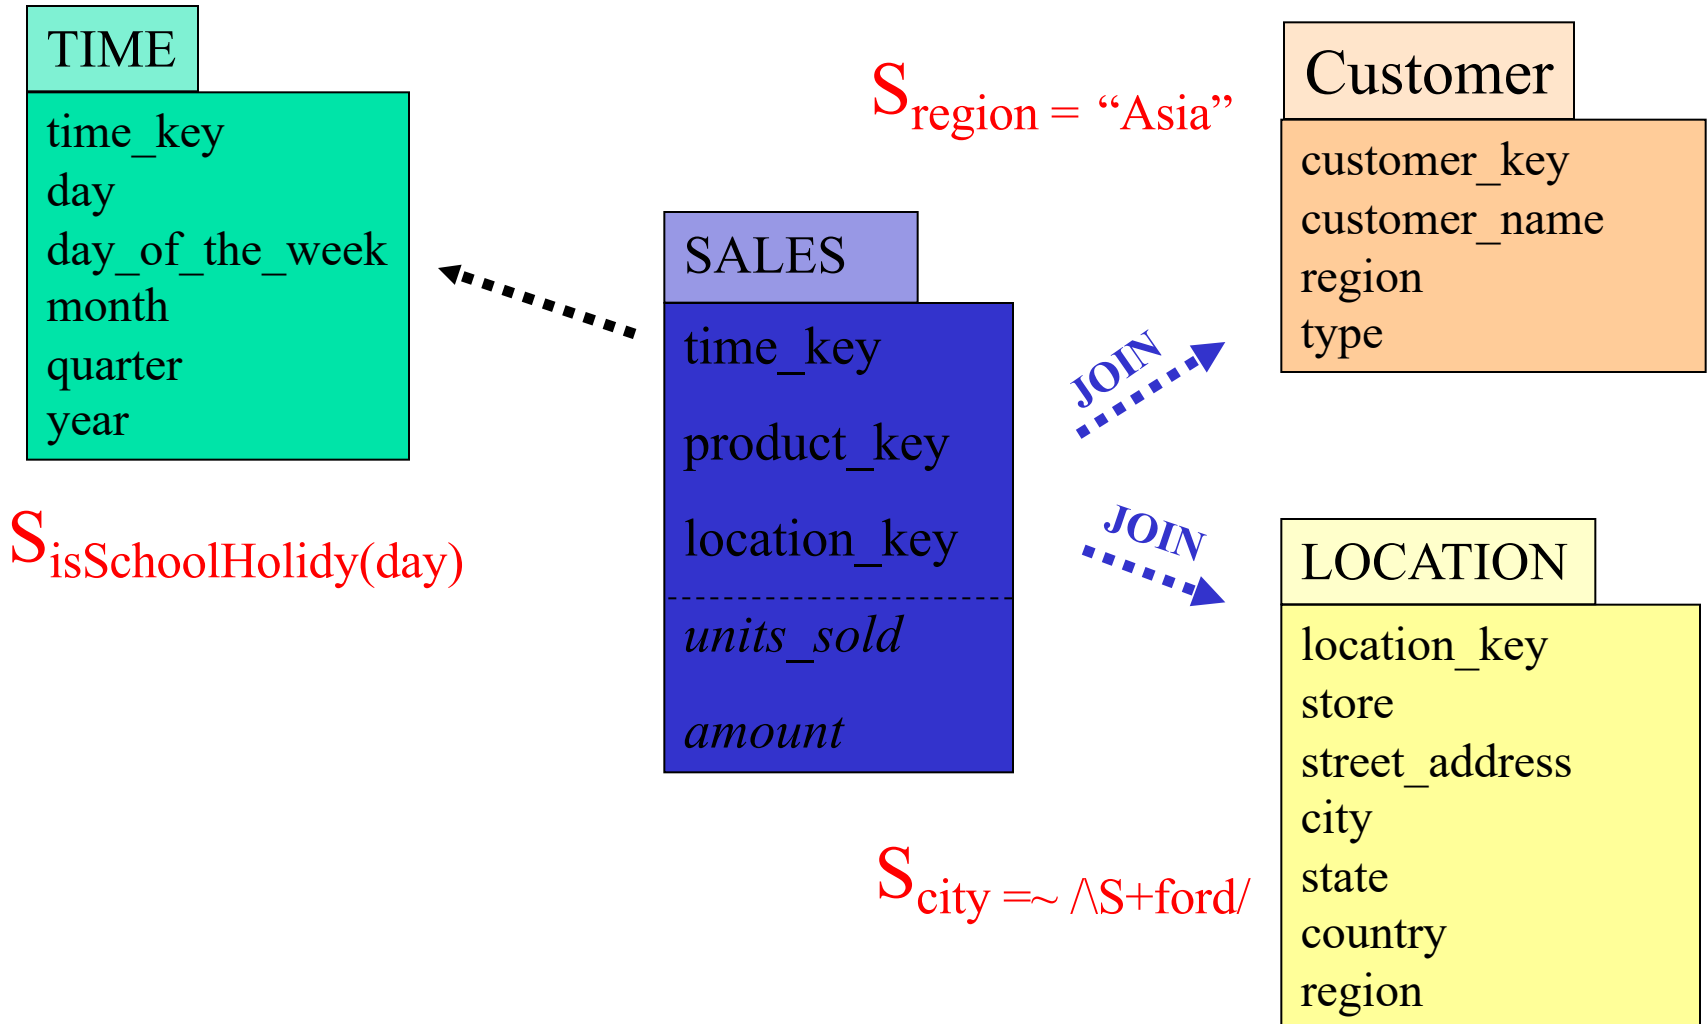


# Indexing OLAP Data: Join Indices

- Join index relates the values of the dimensions of a star schema to rows in the fact table.
  - a join index on *city* maintains for each distinct city a list of ROW-IDs of the tuples recording the sales in the city
- Join indices can span multiple dimensions OR
  - can be implemented as bitmap-indexes (per dimension)
  - use bit-op for multiple-joins



# Q3: Arbitrary selections on Dimensions




# Star Query and Star Join (Cont.)

Usually only part of the dim tables because of the selection predicates

**Time**      **Customer**      **Loc**

time		Cust		loc
1	X	10	X	100
2		20		200



**Sales**

*millions  
of tuples*

time	cust	loc	sold
1	10	100	7
1	10	150	13
1	20	150	2
2	20	200	16
...	...	...	...
1000	2000	500	86

*thousands  
of tuples*

time	cust	loc
1	10	100
1	10	200
1	20	100
1	20	200
2	10	100
2	10	200
2	20	100
2	20	200

$\text{Sales} \triangleright \triangleleft \sigma_1(\text{Time}) \triangleright \triangleleft \sigma_2(\text{Cust}) \triangleright \triangleleft \sigma_3(\text{Loc}) \rightarrow$   
 $\text{Sales} \triangleright \triangleleft (\sigma_1(\text{Time}) \times \sigma_2(\text{Cust}) \times \sigma_3(\text{Loc}))$



# Q4: Coarse-grain Aggregations

- *“Find total sales per customer type in our stores in Europe”*
  - Join-index will prune  $\frac{3}{4}$  of the data (uniform sales), but the remaining  $\frac{1}{4}$  is still large (several millions transactions)
    - Index is unclustered
- High-level aggregations are expensive!!!!

⇒ Long Query Response Times

⇒ Pre-computation is necessary

⇒ Pre-computation is most beneficial

LOCATON

region

country

state

city

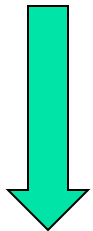
store

# Cuboids = GROUP BYs

---

- Multidimensional aggregation = selection on corresponding cuboid

$GB_{(type, city)}(Sales \triangleright \triangleleft \sigma_1(Time) \triangleright \triangleleft \sigma_2(Cust) \triangleright \triangleleft \sigma_3(Loc))$



$\sigma_1$  selects some *Years*,  $\sigma_2$  selects some *Brands*,  
 $\sigma_3$  selects some *Cities*,

$GB_{(type, city)}(\sigma_{1,2,3}(Cuboid(Year, Type, City)))$

- Materialize some/all of the cuboids
  - A complex decision involving cuboid sizes, query workload, and physical organization

# Two Issues

---

- How to store the materialized cuboids?
- How to compute the cuboids efficiently?

# CUBE BY in ROLAP

Sales		Product				
		1	2	3	4	ALL
Store	1	454	-	-	925	1379
	2	468	800	-	-	1268
	3	296	-	240	-	536
	4	652	-	540	745	1937
	ALL	1870	800	780	1670	5120

4 Group-bys here:  
 (store,product)  
 (store)  
 (product)  
 ()

- Need to write 4 queries!!!
- Compute them independently

Store	Product_key	sum(amtout)
1	1	454
1	4	925
2	1	468
2	2	800
3	1	296
3	3	240
4	1	625
4	3	240
4	4	745
1	ALL	1379
2	ALL	1268
3	ALL	536
4	ALL	1937
ALL	1	1870
ALL	2	800
ALL	3	780
ALL	4	1670
ALL	ALL	5120

**SELECT** LOCATION.store, SALES.product\_key, SUM (amount)

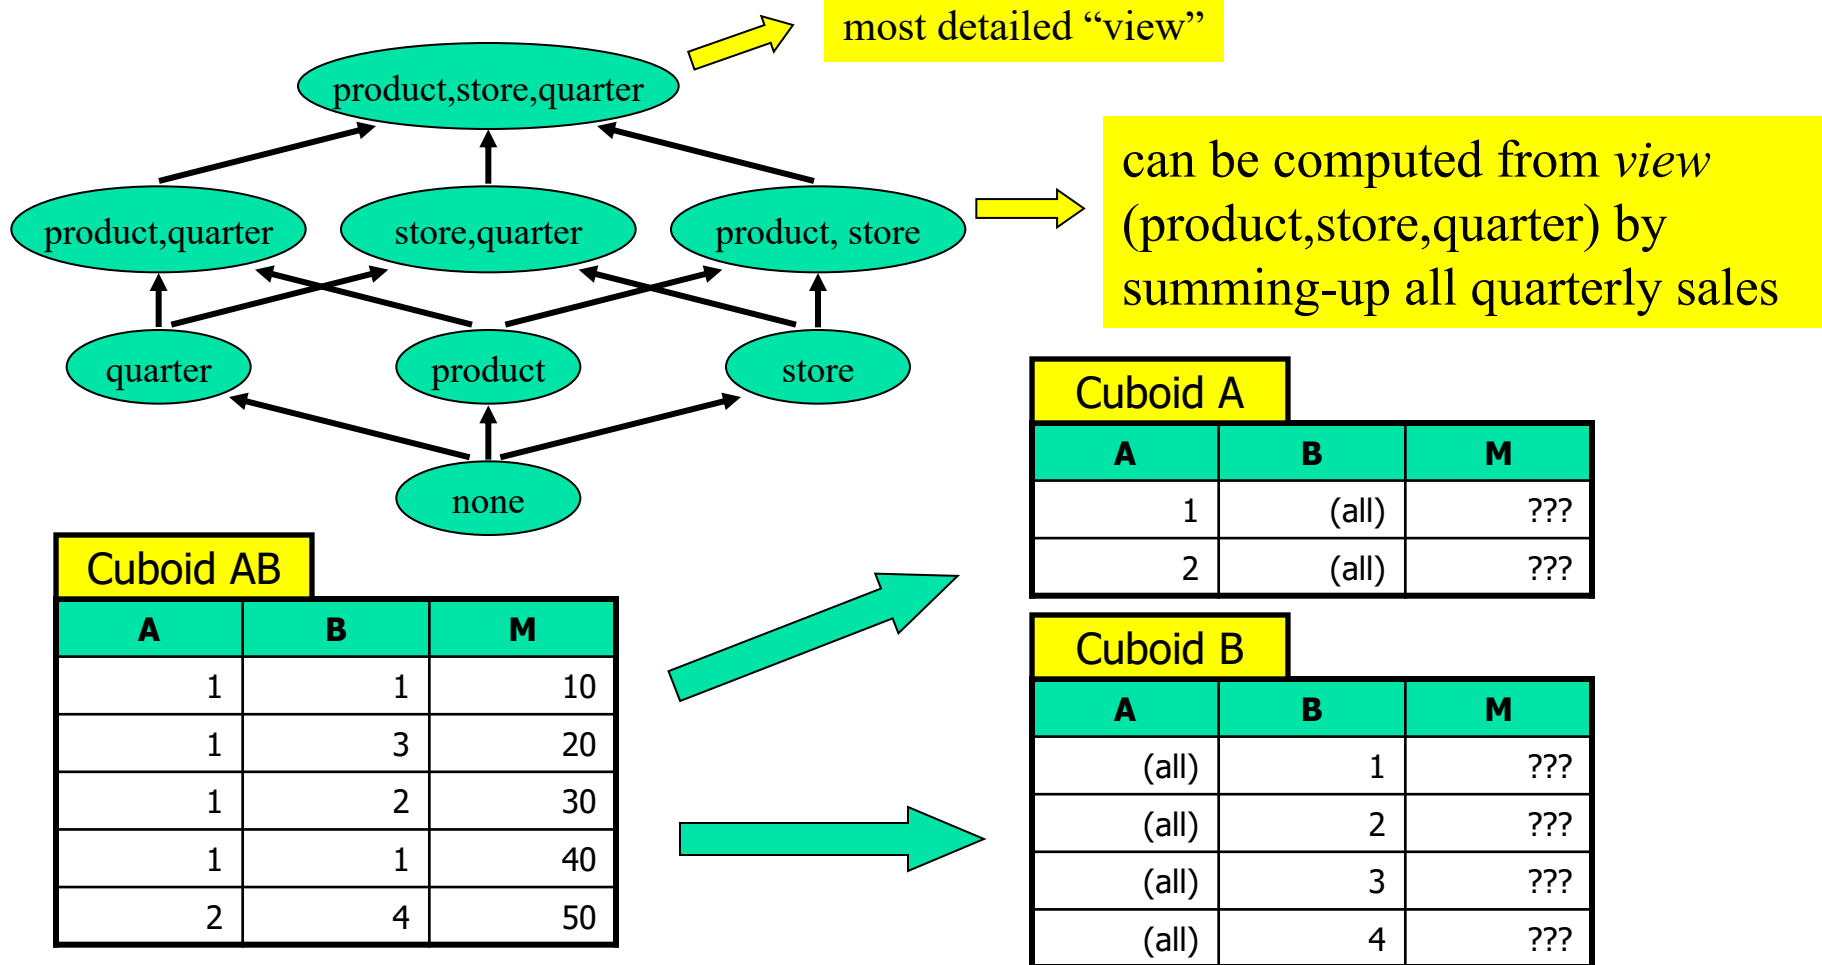
**FROM** SALES, LOCATION

**WHERE** SALES.location\_key=LOCATION.location\_key

**CUBE BY** SALES.product\_key, LOCATION.store

# Top-down Approach


- Model dependencies among the aggregates:




# Bottom-Up Approach (BUC)

- BUC (Beyer & Ramakrishnan, SIGMOD'99)
- Ideas
  - Compute the cube from bottom up
  - Divide-and-conquer
- A simpler recursive version:
  - BUC-SR

A	B	...
1	1	...
1	3	...
1	2	...
1	1	...
2	...	...



■ ■ ■



# Understanding Recursion /1

---

- Powerful computing/problem-solving techniques
- Examples

- Factorial:

- $f(n) = 1$ , if  $n = 1$
- $f(n) = f(n-1) * n$ , if  $n \geq 1$

$$f(0) = 0! =$$
$$???$$

- Quick sort:

- $\text{Sort}([x]) = [x]$
- $\text{Sort}([x_1, \dots, \text{pivot}, \dots, x_n]) = \text{sort}[\text{ys}] ++ \text{sort}[\text{zs}]$ , where


$$\text{ys} = [x \mid x \text{ in } x_i, x \leq \text{pivot}]$$

$$\text{zs} = [x \mid x \leftarrow x_i, x > \text{pivot}]$$

List comprehension  
in Haskell or  
python

# Understanding Recursion /2

- Let  $C(n, m)$  be the number of ways to select  $m$  balls from  $n$  numbered balls
- Show that  $C(n, m) = C(n-1, m-1) + C(n-1, m)$
- Example:  $m = 3, n = 5$
- Consider any ball in the 5 balls, e.g., 'd'





# Key Points

---

- Sub-problems need to be “**smaller**”, so that a simple/trivial boundary case can be reached
- Divide-and-conquer
  - There may be multiple ways the entire solution space can be divided into **disjoint** sub-spaces, each of which can be conquered **recursively**.

# Geometric Intuition /1


- Reduce Cube(in 2D) to Cube(in 1D)

	b1	b2	b3	
a1	M11	M12	M13	[Step 1]
a2	M21	M22	M23	[Step 1]
	[Step 2]	[Step 2]	[Step 2]	[Step 3]

	b1	b2	b3	*
[a1] ×	M11	M12	M13	[Step 1]
[a2] ×	M21	M22	M23	[Step 1]
[*] ×	[Step 2]	[Step 2]	[Step 2]	[Step 3]



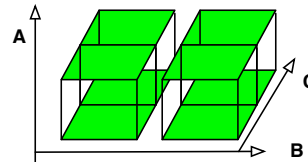
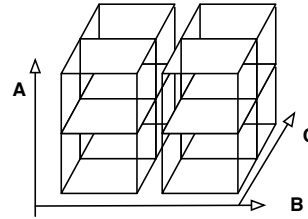
# Geometric Intuition /2

- Reduce Cube(in 3D) to Cube(in 2D)



# Geometric Intuition /3

- Reduce Cube(in 3D) to Cube(in 2D)



# Algebraic Derivation

- How to compute  $n$ -dim cube on  $(n+1)$ -dim base cuboid (array)?
  - What do output tuples look like?
- How to compute  $(n+1)$ -dim cube on  $(n+1)$ -dim base cuboid (array)?
  - What **else** do we need?

[{r1-r5}, ABC]

[{r1-r5}, BC]

	A	B	C	M
r1	1	1	1	10
r2	1	1	2	20
r3	1	2	1	30
r4	1	3	1	40
r5	2	1	1	50

# BUC-SR (Simple Recursion)\*

---

- BUC-SR(data, dims)
  - If (dims is empty)
    - Output (sum(data))
  - Else
    - Dims = [dim1, rest\_of\_dims]
    - For each distinct value  $v$  of dim1
      - slice\_ $v$  = slice of data on “dim1 =  $v$ ”
      - BUC-SR(slice\_ $v$ , rest\_of\_dims)
    - data' = Project(data, rest\_of\_dims)
    - BUC-SR(data', rest\_of\_dims)

Boundary case:  
data is essentially a list  
of measure values

General case:  
1) Slice on dim1. Call  
BUC-SR recursively for  
each slice  
  
2) Project out dim1, and  
call BUC-SR on it  
recursively

# Example

	1	2	3	*
1	30	30	40	100
2	50			50
*	80	30	40	150

$\{\{r1-r5\}, AB\}$

A	B	M
1	1	10
1	1	20
1	2	30
1	3	40
2	1	50

r1  
r2  
r3  
r4  
r5

$\{\{r1-r4\}, B\}$

B	M
1	10
1	20
2	30
3	40

$\{\{r5\}, B\}$

B	M
1	50

$\{\{r1'-r5'\}, B\}$

B	M
1	80
2	30
3	40

Internal Output

B	M
1	30
2	30
3	40
*	100

B	M
1	50
*	50

B	M
1	80
2	30
3	40
*	150

Output

A	B	M
1	1	30
1	2	30
1	3	40
1	*	100

A	B	M
2	1	50
2	*	50

A	B	M
*	1	80
*	2	30
*	3	40
*	*	150

# Try a 3D-Cube by Yourself

---

[{r1-r5}, ABC]

	A	B	C	M
r1	1	1	1	10
r2	1	1	2	20
r3	1	2	1	30
r4	1	3	1	40
r5	2	1	1	50



# MOLAP

---

- (Sparse) array-based multidimensional storage engine
- Pros:
  - small size (esp. for dense cubes)
  - fast in indexing and query processing
- Cons:
  - scalability
  - conversion from relational data

# Multidimensional Array

$$f(\text{time}, \text{item}) = 4 * \text{time} + \text{item}$$

time	item	dollars_sold
Q1	home entertainment	605
Q2	home entertainment	680
Q3	home entertainment	812
Q4	home entertainment	927
Q1	computer	825
Q2	computer	952
Q3	computer	1023
Q4	computer	1038
Q1	phone	14
Q2	phone	31
Q3	phone	30
Q4	phone	38
Q1	security	400
Q2	security	512
Q3	security	501
Q4	security	580



Step 1

Mappings

time	value
Q1	0
Q2	1
Q3	2
Q4	3

item	value
home entertainment	0
computer	1
phone	2
security	3

time	item	dollars_sold
0	0	605
1	0	680
2	0	812
3	0	927
0	1	825
1	1	952
2	1	1023
3	1	1038
0	2	14
1	2	31
2	2	30
3	2	38
0	3	400
1	3	512
2	3	501
3	3	580


offset

0  
4  
8  
12  
1  
5  
9  
13  
2  
6  
10  
14  
3  
7  
11  
15

# Multidimensional Array

Step 3: If **dense**, only need to store sorted slots

offset	dollars_sold
0	605
1	825
2	14
3	400
4	680
5	952
6	31
7	512
8	812
9	1023
10	30
11	501
12	927
13	1038
14	38
15	580



- Think: how to decode a slot?

Dense MD array
605
825
14
400
680
952
31
512
812
1023
30
501
927
1038
38
580

$$f(\text{time}, \text{item}) = 4 * \text{time} + \text{item}$$

# The Sparse Case

time	item	dollars_sold
Q1	home entertainment	605
/* same table but with many rows deleted to make it sparse */		
Q4	security	580



Step 1

Mappings

time	value
Q1	0
Q2	1
Q3	2
Q4	3

item	value
home entertainment	0
computer	1
phone	2
security	3

time	item	dollars_sold	offset
0	0	605	0
/* same table but with many rows deleted to make it sparse */			
3	3	580	15

# Multidimensional Array

## Choice 1

offset	dollars_sold
0	605
15	580

- Think: how to decode a slot?
- Multidimensional array is typically sparse
  - Use sparse array (i.e., offset + value)
  - Could use chunk to further reduce the space
- Space usage:
  - $(d+1)*n*4$  vs  $2*n*4$
- HOLAP:
  - Store all non-base cuboid in MD array
  - Assign a value for ALL

## Choice 2

Dense MD array
605
-
-
-
-
-
-
-
-
-
-
-
-
-
580