


# Approximate Nearest Neighbor Search based on Product Quantization

Wei Wang

CSE, UNSW

# Nearest Neighbor (NN) and its Variants



- Given:  $n$  points in  $d$ -dimensional Euclidean space
- NN query: find the closest point,  $o^*$ , in  $D$ , wrt a query point  $q$
- Approximate NN query: find a near point
- All extendible to  $k$ NN versions for  $k > 1$

# Motivations

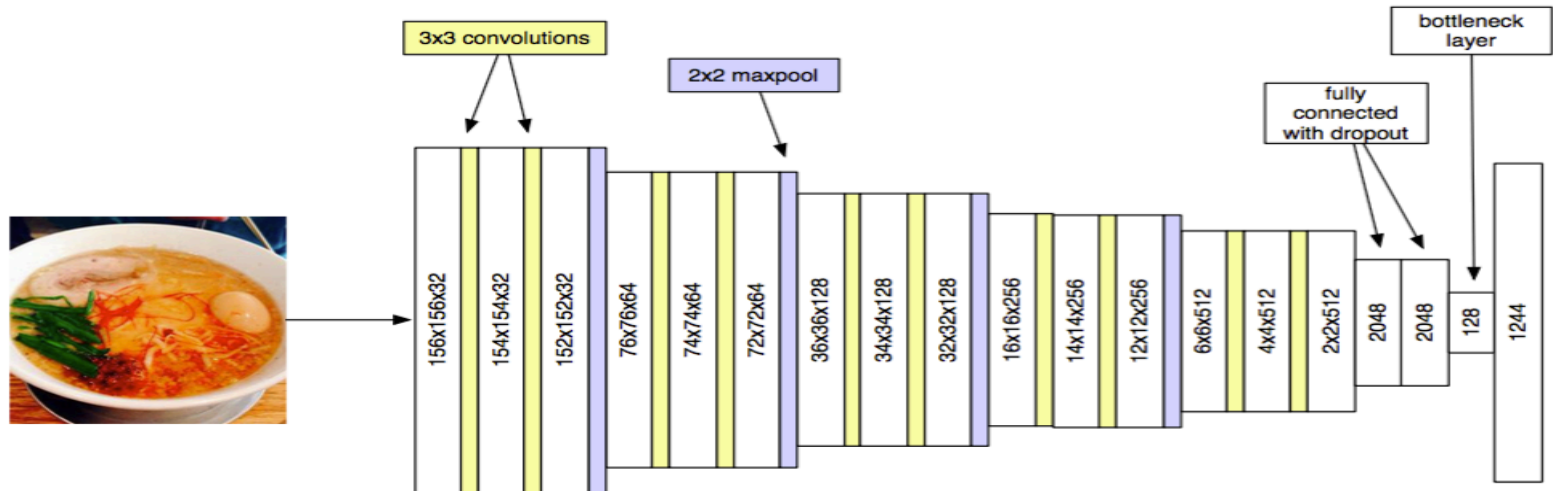
- Theoretical interest
  - Fundamental geometric problem: “post-office problem”
  - Known to be hard due to high-dimensionality
- Many applications
  - Feature vectors: Data Mining, Multimedia DB
  - Applications based on similarity queries, e.g.,
    - Quantization in coding/compression
    - Recommendation systems
    - Bioinformatics/Chemoinformatics
  - Machine Learning
    - Representation learning (e.g., embedding)
    - Dimensionality reduction; collaborative filtering; kNN classifier; Kernel?

# Dimensionality reduction for ML

- Start with high-dimensional data
- Run dimensionality reduction
- Do stuff in a small dimensional space

# Deep Learning for food

- Deep model trained on a GPU on 6M random pics downloaded from Yelp



# Distance in Smaller Space

1. Run image through the network
2. Use the 128-dimensional bottleneck layer as an item vector
3. Use cosine distance in the reduced space

$$f\left(\text{img}\right) = \mathbf{v} \quad (\mathbf{v} \in \mathbb{R}^{128})$$

$$d(\mathbf{u}, \mathbf{v}) = \left( \frac{\mathbf{u}}{|\mathbf{u}|} - \frac{\mathbf{v}}{|\mathbf{v}|} \right)^2$$

# Nearest Neighbor Food Pics



# Preliminary: Vector Quantization (with K-means)


- Idea: compressed representation of vectors
  - Each D-dim vector  $x$  is represented by  $QZ(x)$ , where  $QZ()$  is a quantizer
  - $\text{dist}(x, o) \approx \text{dist}(x, QZ(o))$
- Encode the vectors:
  - Learn a codebook  $W = \{c_1, c_2, \dots, c_K\}$  via K-means
  - Assign  $o$  to its **nearest codeword** in  $W$ 
    - E.g.,  $QZ(o) = c_i$  ( $i \in 1 \dots K$ ), such that  $\text{dist}(x, c_i) \leq \text{dist}(x, c_j)$  for  $\forall j$
  - Represent each vector  $o$  by its assigned codeword
- Assume  $D = 256$ ,  $K = 2^{16}$ 
  - Before: 4 bytes \* 256 = 1024 bytes for each vector
  - Now:
    - data: 16 bits = 2 bytes
    - codebook:  $4 * 256 * 2^{16}$

smaller when  $n > 4 * 256 * 2^{16} / 1022 = 65,664$




# Vector Quantization – Query Processing

- Given query  $q$ , how to find a point close to  $q$ ?



# Vector Quantization – Query Processing

- Given query  $q$ , how to find a point close to  $q$ ?
- Algorithm:
  1. Compute  $QZ(q)$
  2. **Candidate set  $C$**  = all data vectors associated with  $QZ(q)$
  3. Verification: compute distance between  $Q$  and  $o_i \in C$ 
    - Requires loading the vectors in  $C$




- Any problem/improvement?

**Inverted index:** a hash table that maps  $c_j$  to a list of  $o_i$  that are associated with  $c_j$

# Limitations of VQ

- To achieve better accuracy, fine-grained quantizer with large  $K$  is needed
- Large  $K \rightarrow$ 
  - Costly to run K-means
  - Computing  $QZ(q)$  is expensive:  $O(K M)$
  - May need to look beyond  $q(Q)$  cell
- Solutions:
  - **Product Quantization**, or
  - Hierarchical k-means



# Product Quantization

- Idea:
  - Partition the dimension into  $m$  partitions
    - Accordingly a vector  $\rightarrow m$  subvectors
  - Use separate VQ with  $k$  codewords for each chunk

- Example:

- 10-dim vector decomposed in  $m = 2$  subvectors
  - $\mathbf{o}^T = [\mathbf{o}_1 : \mathbf{o}_2]^T$
- Each codebook has 4 codewords, denoted as  $c_{i,j}$
- Total space in bits:
  - data:  $m \log(K)$
  - codebook:  $m * (D/m) * K$

7	2	-1	5	6	4	3	-5	-1	-7
-2	3	0	6	4	2	2	4	5	-8
...	...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...	...

0	1	0	0
1	0	1	1
...	...	...	...
...	...	...	...
...	...	...	...

$C_{1,1}$	5.9	2.3	-2.7	3.9	6.1
$C_{1,2}$	-1.3	1.8	7.4	5.5	0.9
$C_{1,3}$	...	...	...	...	...
$C_{1,4}$	...	...	...	...	...

$C_{2,1}$	-0.1	3.5	1.4	9.6	5.5
$C_{2,2}$	...	...	...	...	...
$C_{2,3}$	...	...	...	...	...
$C_{2,4}$	...	...	...	...	...

# Distance Estimation

- Euclidean distance between a query point **q** and a data point encoded as **t**
  - Restore the virtual joint center by looking up each partition of **t** in the corresponding codebooks  $\rightarrow$  **p**

$$d^2(\mathbf{q}, \mathbf{t}) = \sum_{i=1}^D (\mathbf{q}_i - \mathbf{p}_i)^2$$

- This is the **Asymmetric** Distance (AD)

$$d^2(\mathbf{q}, \mathbf{t}) = \sum_{i=1}^m (\mathbf{q}_{(i)} - \mathbf{c}_{i, \mathbf{t}_{(i)}})^2$$

**q**

7	2	-1	5	6	4	3	-5	-1	-7
---	---	----	---	---	---	---	----	----	----

**t**

0	1	0	0
---	---	---	---



$C_{1,1}$	5.9	2.3	-2.7	3.9	6.1
$C_{1,2}$	-1.3	1.8	7.4	5.5	0.9
$C_{1,3}$	...	...	...	...	...
$C_{1,4}$	...	...	...	...	...

$C_{2,1}$	-0.1	3.5	1.4	9.6	5.5
$C_{2,2}$	...	...	...	...	...
$C_{2,3}$	...	...	...	...	...
$C_{2,4}$	...	...	...	...	...

**p**

-1.3	1.8	7.4	5.5	0.9	-0.1	3.5	1.4	9.6	5.5
------	-----	-----	-----	-----	------	-----	-----	-----	-----

# PQ – Query Processing

- Naïve:
  - Perform ADC for every  $\mathbf{t}$  in the database
  - Candidate = those with the  $k$  smallest AD
  - [Optional] Reranking (if  $k > 1$ ):
    - Load the data vectors and compute the actual Euclidean distance
    - Return the one with the smallest distance
- Pruning:
  - AD is monotonic in each component

$$d^2(\mathbf{q}, \mathbf{t}) = \sum_{i=1}^m (\mathbf{q}_{(i)} - \mathbf{c}_{i, \mathbf{t}_{(i)}})^2$$

There is an optimal order to consider each encoded points!

Compute it efficiently and incrementally!

$\mathbf{q}$

7	2	-1	5	6	4	3	-5	-1	-7
---	---	----	---	---	---	---	----	----	----

$\mathbf{t}$

$\mathbf{t}'$

0	1	0	0
1	0	1	1
1	0	1	0
0	0	1	0
1	1	1	1

$$(\mathbf{q}_{(1)} - \mathbf{c}_{i, \mathbf{t}_{(1)}})^2 + (\mathbf{q}_{(2)} - \mathbf{c}_{i, \mathbf{t}_{(2)}})^2$$

$$(\mathbf{q}_{(1)} - \mathbf{c}_{i, \mathbf{t}'_{(1)}})^2 + (\mathbf{q}_{(2)} - \mathbf{c}_{i, \mathbf{t}'_{(2)}})^2$$

# Multi-index Algorithm


1. Sort each code book in increasing order of their partial AD distance to  $\mathbf{q}$

2. Maintain the non-dominated combinations into a min heap  $H$

NOTE: slightly different from the paper – Skyline algorithm

$\mathbf{q}_{(1)} \text{ vs. } \mathcal{U}$			$\mathbf{q}_{(2)} \text{ vs. } \mathcal{V}$		
$i$	$\mathbf{u}_{\alpha(i)}$	$r$	$j$	$\mathbf{v}_{\beta(j)}$	$s$
1	$\mathbf{u}_3$	0.5	1	$\mathbf{v}_4$	0.1
2	$\mathbf{u}_4$	0.7	2	$\mathbf{v}_3$	2
3	$\mathbf{u}_5$	4	3	$\mathbf{v}_5$	3
4	$\mathbf{u}_2$	6	4	$\mathbf{v}_2$	6
5	$\mathbf{u}_1$	8	5	$\mathbf{v}_6$	7
6	$\mathbf{u}_6$	9	6	$\mathbf{v}_1$	11


partial AD



 dominates  dominates 

$(i,j) \rightarrow$  pushes  $(i+1, j)$  and  $(i, j+1)$  into  $H$

Dedup: need to check if it has EVER been pushed into  $H$  or not



NOTE: Reusing figure in the paper.  $m = 2$ , each code book has 6 codewords, named  $\mathbf{u}_i$  and  $\mathbf{v}_j$

# Multi-index Algorithm

1. Sort each code book in increasing order of their partial AD distance to  $\mathbf{q}$
2. Maintain the non-dominated combinations into a min heap  $H$


$\mathbf{q}_{(1)} \text{ vs. } \mathcal{U}$

$i$	$\mathbf{u}_{\alpha(i)}$	$r$
1	$\mathbf{u}_3$	0.5
2	$\mathbf{u}_4$	0.7
3	$\mathbf{u}_5$	4
4	$\mathbf{u}_2$	6
5	$\mathbf{u}_1$	8
6	$\mathbf{u}_6$	9

$\mathbf{q}_{(2)} \text{ vs. } \mathcal{V}$

$j$	$\mathbf{v}_{\beta(j)}$	$s$
1	$\mathbf{v}_4$	0.1
2	$\mathbf{v}_3$	2
3	$\mathbf{v}_5$	3
4	$\mathbf{v}_2$	6
5	$\mathbf{v}_6$	7
6	$\mathbf{v}_1$	11

partial AD



(2,2) not pushed again

(2,3) not pushed again

## Exercises:

- Write out  $H$
- How to perform dedup?
- Generalize to  $m$ -dimensional case
- [hard] perform dedup without additional data structure (and cost)



# References

- Hervé Jégou, Matthijs Douze, and Cordelia Schmid. Product quantization for nearest neighbor search. TPAMI 2014.
- Artem Babenko and Victor Lempitsky. The Inverted Multi-Index. TPAMI 2014.